

Amazon Fine Food Reviews Analysis

Data Source: <https://www.kaggle.com/snap/amazon-fine-food-reviews> (<https://www.kaggle.com/snap/amazon-fine-food-reviews>)

The Amazon Fine Food Reviews dataset consists of reviews of fine foods from Amazon.

Number of reviews: 568,454 Number of users: 256,059 Number of products: 74,258 Timespan: Oct 1999 - Oct 2012 Number of Attributes/Columns in data: 10

Attribute Information:

- 1.Id
- 2.ProductId - unique identifier for the product
- 3.UserId - unique identifier for the user
- 4.ProfileName
- 5.HelpfulnessNumerator - number of users who found the review helpful
- 6.HelpfulnessDenominator - number of users who indicated whether they found the review helpful or not
- 7.Score - rating between 1 and 5
- 8.Time - timestamp for the review
- 9.Summary - brief summary of the review
- 10.Text - text of the review

Objective:

Given a review, determine whether the review is positive (rating of 4 or 5) or negative (rating of 1 or 2).

1. Import required libraries

```
In [2]: 1 import warnings
        2 warnings.filterwarnings("ignore")
```

```
In [3]: 1 %matplotlib inline
2
3 import sqlite3
4 import pandas as pd
5 import numpy as np
6 import nltk
7 import string
8 import matplotlib.pyplot as plt
9 import seaborn as sns
10 from sklearn.feature_extraction.text import TfidfTransformer
11 from sklearn.feature_extraction.text import TfidfVectorizer
12
13 from sklearn.feature_extraction.text import CountVectorizer
14 from sklearn import metrics
15 from sklearn.model_selection import train_test_split
16 import re
17 # Tutorial about Python regular expressions: https://pymotw.com/2/re/
18 import string
19 from nltk.corpus import stopwords
20 from nltk.stem import SnowballStemmer
21 from nltk.stem.wordnet import WordNetLemmatizer
22 from gensim.models import Word2Vec
23 from gensim.models import KeyedVectors
24 import pickle
25
26 from tqdm import tqdm_notebook
27 from tqdm import tqdm
28 from bs4 import BeautifulSoup
29 import os
30 from keras.models import Sequential
31 from keras.layers import Dense
32 from keras.layers import LSTM
33 from keras.layers.embeddings import Embedding
34 from keras.preprocessing import sequence
35 from keras.preprocessing.text import Tokenizer
36
37 from prettytable import PrettyTable
38 # fix random seed for reproducibility
39 np.random.seed(7)
```

Using TensorFlow backend.

2. Read the Dataset

- a. Create a Connection object that represents the database. Here the data will be stored in the 'database.sqlite' file.
- b. Read the Dataset table using connection object where the score column != 3
- c. Replace the score values with 'positive' and 'negative' label.(i.e Score 1 & 2 is labeled as negative and Score 4 & 5 is labeled as positive)
- d. Score with value 3 is neutral.

```

In [3]: 1 # using SQLite Table to read data.
2 con = sqlite3.connect('database.sqlite')
3
4 # filtering only positive and negative reviews i.e.
5 # not taking into consideration those reviews with Score=3
6 # SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000, will give top 500000 data points
7 # you can change the number to any other number based on your computing power
8
9 # filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 500000""", con)
10 # for tsne assignment you can take 5k data points
11
12 filtered_data = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 LIMIT 200000""", con)
13
14 # Give reviews with Score>3 a positive rating(1), and reviews with a score<3 a negative rating(0).
15 def partition(x):
16     if x < 3:
17         return 0
18     return 1
19
20 #changing reviews with score less than 3 to be positive and vice-versa
21 actualScore = filtered_data['Score']
22 positiveNegative = actualScore.map(partition)
23 filtered_data['Score'] = positiveNegative
24 print("Number of data points in our data", filtered_data.shape)
25 filtered_data.head(3)

```

Number of data points in our data (200000, 10)

Out[3]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
0	1	B001E4KFG0	A3SGXH7AUHU8GW	delmartian	1	1	1	1303862400	Good Quality Dog Food	I have bought several of the Vitality canned d...

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
1	2	B00813GRG4	A1D87F6ZCVE5NK	dll pa	0	0	0	1346976000	Not as Advertised	Product arrived labeled as Jumbo Salted Peanut...
2	3	B000LQOCH0	ABXLMWJIXXAIN	Natalia Corres "Natalia Corres"	1	1	1	1219017600	"Delight" says it all	This is a confection that has been around a fe...

```
In [4]: 1 display = pd.read_sql_query("""
2 SELECT UserId, ProductId, ProfileName, Time, Score, Text, COUNT(*)
3 FROM Reviews
4 GROUP BY UserId
5 HAVING COUNT(*)>1
6 """, con)
```

```
In [5]: 1 print(display.shape)
2 display.head()
```

(80668, 7)

Out[5]:

		UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
0	#oc-R115TNMSPFT9I7	B007Y59HVM	Breyton	1331510400	2	Overall its just OK when considering the price...		2
1	#oc-R11D9D7SHXIJB9	B005HG9ET0	Louis E. Emory "hoppy"	1342396800	5	My wife has recurring extreme muscle spasms, u...		3
2	#oc-R11DNU2NBKQ23Z	B007Y59HVM	Kim Cieszykowski	1348531200	1	This coffee is horrible and unfortunately not ...		2
3	#oc-R11O5J5ZVQE25C	B005HG9ET0	Penguin Chick	1346889600	5	This will be the bottle that you grab from the...		3
4	#oc-R12KPBODL2B5ZD	B007OSBE1U	Christopher P. Presta	1348617600	1	I didnt like this coffee. Instead of telling y...		2

```
In [6]: 1 display[display['UserId']=='AZY10LLTJ71NX']
```

Out[6]:

	UserId	ProductId	ProfileName	Time	Score	Text	COUNT(*)
80638	AZY10LLTJ71NX	B006P7E5ZI	undertheshrine "undertheshrine"	1334707200	5	I was recommended to try green tea extract to ...	5

```
In [7]: 1 display['COUNT(*)'].sum()
```

Out[7]: 393063

4. Exploratory Data Analysis

Data Cleaning: Deduplication

It is observed (as shown in the table below) that the reviews data had many duplicate entries. Hence it was necessary to remove duplicates in order to get unbiased results for the analysis of the data. Following is an example:

```
In [8]: 1 display= pd.read_sql_query("""
2 SELECT *
3 FROM Reviews
4 WHERE Score != 3 AND UserId="AR5J8UI46CURR"
5 ORDER BY ProductID
6 """, con)
7 display.head()
```

Out[8]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
0	78445	B000HDL1RQ	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS	DELICIOUS WAFERS FIND THEM EUROPEAN WAFERS
1	138317	B000HDOPYC	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS	DELICIOUS WAFERS FIND THEM EUROPEAN WAFERS
2	138277	B000HDOPYM	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS	DELICIOUS WAFERS FIND THEM EUROPEAN WAFERS
3	73791	B000HDOPZG	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS	DELICIOUS WAFERS FIND THEM EUROPEAN WAFERS
4	155049	B000PAQ75C	AR5J8UI46CURR	Geetha Krishnan	2	2	5	1199577600	LOACKER QUADRATINI VANILLA WAFERS	DELICIOUS WAFERS FIND THEM EUROPEAN WAFERS

As it can be seen above that same user has multiple reviews with same values for HelpfulnessNumerator, HelpfulnessDenominator, Score, Time, Summary and Text and on doing analysis it was found that

It was inferred after analysis that reviews with same parameters other than ProductId belonged to the same product just having different flavour or quantity. Hence in order to reduce redundancy it was decided to eliminate the rows having same parameters.

The method used for the same was that we first sort the data according to ProductId and then just keep the first similar product review and delete the others. for eg. in the above just the review for ProductId=B000HDL1RQ remains. This method ensures that there is only one representative for each product and deduplication without sorting would lead to possibility of different representatives still existing for the same product.

```
In [9]: 1 #Sorting data according to ProductId in ascending order
        2 sorted_data=filtered_data.sort_values('ProductId', axis=0, ascending=True, inplace=False, kind='quicksort', na_posit
```

```
In [10]: 1 #Deduplication of entries
        2 final=sorted_data.drop_duplicates(subset={"UserId","ProfileName","Time","Text"}, keep='first', inplace=False)
        3 final.shape
```

Out[10]: (160178, 10)

```
In [11]: 1 #Checking to see how much % of data still remains
        2 (final['Id'].size*1.0)/(filtered_data['Id'].size*1.0)*100
```

Out[11]: 80.089

Observation:- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed from calculations


```
In [12]: 1 display= pd.read_sql_query("""
2         SELECT *
3         FROM Reviews
4         WHERE Score != 3 AND Id=44737 OR Id=64422
5         ORDER BY ProductID
6         """, con)
7
8         display.head(2)
```

Out[12]:

	Id	ProductId	UserId	ProfileName	HelpfulnessNumerator	HelpfulnessDenominator	Score	Time	Summary	Text
0	64422	B000MIDROQ	A161DK06JJMCYF	J. E. Stephens "Jeanne"	3	1	5	1224892800	Bought This for My Son at College	My son loves spaghetti so I didn't hesitate or...
1	44737	B001EQ55RW	A2V0I904FH7ABY	Ram	3	2	4	1212883200	Pure cocoa taste with crunchy almonds inside	It was almost a 'love at first bite' - the per...

- It was also seen that in two rows given below the value of HelpfulnessNumerator is greater than HelpfulnessDenominator which is not practically possible hence these two rows too are removed

```
In [13]: 1 final=final[final.HelpfulnessNumerator<=final.HelpfulnessDenominator]
```

```
In [14]: 1 #Before starting the next phase of preprocessing Lets see the number of entries left
          2 print(final.shape)
          3
          4 #How many positive and negative reviews are present in our dataset?
          5 final['Score'].value_counts()
```

(160176, 10)

```
Out[14]: 1    134799
          0     25377
          Name: Score, dtype: int64
```

5. Preprocessing

[5.1]. Preprocessing Review Text and Summary

Now that we have finished deduplication our data requires some preprocessing before we go on further with analysis and making the prediction model.

Hence in the Preprocessing phase we do the following in the order below:-

1. Begin by removing the html tags
2. Remove any punctuations or limited set of special characters like , or . or # etc.
3. Check if the word is made up of english letters and is not alpha-numeric
4. Check to see if the length of the word is greater than 2 (as it was researched that there is no adjective in 2-letters)
5. Convert the word to lowercase
6. Remove Stopwords
7. Finally Snowball Stemming the word (it was observed to be better than Porter Stemming)

After which we collect the words used to describe positive and negative reviews

In [15]:

```

1  # https://stackoverflow.com/a/47091490/4084039
2  import re
3
4  def decontracted(phrase):
5      # specific
6      phrase = re.sub(r"won't", "will not", phrase)
7      phrase = re.sub(r"can't", "can not", phrase)
8
9      # general
10     phrase = re.sub(r"n't", " not", phrase)
11     phrase = re.sub(r"\ 're", " are", phrase)
12     phrase = re.sub(r"\ 's", " is", phrase)
13     phrase = re.sub(r"\ 'd", " would", phrase)
14     phrase = re.sub(r"\ 'll", " will", phrase)
15     phrase = re.sub(r"\ 't", " not", phrase)
16     phrase = re.sub(r"\ 've", " have", phrase)
17     phrase = re.sub(r"\ 'm", " am", phrase)
18     return phrase

```

In [16]:

```

1  # Combining all the above students
2  from tqdm import tqdm
3  def createCleanedText(review_text, column_name):
4      sno = nltk.stem.SnowballStemmer('english') #initialising the snowball stemmer
5      preprocessed_reviews = []
6      # tqdm is for printing the status bar
7      for sentence in tqdm(review_text):
8          sentence = re.sub(r"http\S+", "", sentence) # \S=except space; + = 1 or more
9          sentence = BeautifulSoup(sentence, 'lxml').get_text() # remove links
10         sentence = decontracted(sentence) # expand short forms
11         sentence = re.sub("\S*\d\S*", "", sentence).strip() #remove words containing digits
12         sentence = re.sub('[^A-Za-z]+', ' ', sentence) # remove special char
13         # https://gist.github.com/sebleier/554280
14         sentence = ' '.join(e.lower() for e in sentence.split())
15         preprocessed_reviews.append(sentence.strip())
16         #adding a column of CleanedText which displays the data after pre-processing of the review
17         final[column_name]=preprocessed_reviews
18

```

```
In [17]: 1 if not os.path.isfile('final.sqlite'):
2         #createCleanedText(final['Text_Summary'].values,column_name='CleanedTextSumm')
3         createCleanedText(final['Text'].values,column_name='CleanedText')
4         conn = sqlite3.connect('final.sqlite')
5         c=conn.cursor()
6         conn.text_factory = str
7         final.to_sql('Reviews', conn, schema=None, if_exists='replace', \
8                     index=True, index_label=None, chunksize=None, dtype=None)
9         conn.close()
```

100%|██████████| 160176/160176 [00:54<00:00, 2915.94it/s]

```
In [18]: 1 if os.path.isfile('final.sqlite'):
2         conn = sqlite3.connect('final.sqlite')
3         final = pd.read_sql_query(""" SELECT * FROM Reviews WHERE Score != 3 """, conn)
4         conn.close()
5     else:
6         print("Please the above cell")
```

In [19]:

```
1 print(final.head(3))
2 final.shape
```

```

      index      Id  ProductId      UserId  ProfileName  \
0  138695  150513  0006641040  ASH0DZQQF6AIZ    tessarat
1  138707  150525  0006641040  A2QID6VCFTY51R         Rick
2  138708  150526  0006641040  A3E9QZFE9KXH8J  R. Mitchell

      HelpfulnessNumerator  HelpfulnessDenominator  Score      Time  \
0                        0                        0      1  1325721600
1                        1                        2      1  1025481600
2                       11                       18      0  1129507200

                                Summary  \
0                                A classic
1  In December it will be, my snowman's anniversa...
2                                awesome book poor size

                                Text  \
0  I remembered this book from my childhood and g...
1  My daughter loves all the "Really Rosie" books...
2  This is one of the best children's books ever ...

                                CleanedText
0  i remembered this book from my childhood and g...
1  my daughter loves all the really rosie books s...
2  this is one of the best children is books ever...
```

Out[19]: (160176, 12)

6. Splitting data into Train and Test set

Function for saving and loading current state of the model:

```
In [4]: 1 #Functions to save objects for later use and retireve it
2 def savetofile(obj,filename):
3     pickle.dump(obj,open(filename+".pkl","wb"))
4 def openfromfile(filename):
5     temp = pickle.load(open(filename+".pkl","rb"))
6     return temp
```

```
In [10]: 1 #TEXT COLUMN
2 if os.path.isfile('X'):
3     X=openfromfile('X')
4     y=openfromfile('y')
5     x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=.3)
6 else:
7     X=np.array(final['CleanedText'])
8     y=np.array(final['Score'])
9     savetofile(X,'X')
10    savetofile(y,'y')
11    x_train,x_test,y_train,y_test=train_test_split(X,y,test_size=.3)
```

Function for training model and draw loss plot:

```
In [13]: 1 def trainModel(model, X_train, Y_train, optimizer, loss, epoch, batch_size):
2     #COMPILE MODEL
3     model.compile(optimizer=optimizer,metrics=['accuracy'],loss=loss)
4     #FIT MODEL ON TRAIN DATA AND MEASURE BOTH TRAIN AND CV SCORE
5     history=model.fit(X_train,Y_train,batch_size=batch_size,epochs=epoch,verbose=1,validation_split=.3)
6
7     #PLOT EPOCH VS LOSS
8     plt.figure(1,figsize=(8,5))
9     sns.set_style("darkgrid")
10    plt.title('Epoch vs Loss')
11    plt.plot(range(epoch), history.history['val_loss'], 'r',label='Validation Loss')
12    plt.plot(range(epoch), history.history['loss'], 'b',label='Train Loss')
13    plt.xlabel('No. of Epoch')
14    plt.ylabel('Loss')
15    plt.legend()
16    plt.show()
17    return model, history
```

Function for model performance on test data:

```
In [27]: 1 def model_performance(X_test, Y_test, epoch, history, model_arch):
2         score=model.evaluate(X_test, Y_test)
3         print('Test score: %.4f'%score[0])
4         print('Test accuracy: %.4f'%score[1])
5
6         local_summary=PrettyTable()
7         local_summary.field_names = ["Model", \
8                                     "Test-Loss", "Test-Accuracy",\
9                                     "Train-Loss", "Train-Accuracy", "Val-Loss", "Val-Accu"]
10        local_summary.add_row([model_arch, '%.4f' %score[0], '%.4f' %score[1],\
11                               '%.4f' %history.history['loss'][epoch-1], '%.4f' %history.history['acc'][epoch-1],\
12                               '%.4f' %history.history['val_loss'][epoch-1], '%.4f' %history.history['val_acc'][epoch-1]])
13        return local_summary
14
```

Function for saving model for future use:

```
In [36]: 1 def saveModel(model_obj,model_name):
2         model_obj.save(model_name+'.h5')
```

Initialize common objects :

```
In [15]: 1 vocab_size=8000
2         nb_epoch=5
3         batch_size=128
4         optimizer='adam'
5         loss='binary_crossentropy'
6         embedding_vecor_length = 32
7         max_review_length = 600
```

Preparing text data for model:

```
In [16]: 1 #REFERENCE LINK: https://machinelearningmastery.com/prepare-text-data-deep-learning-keras/
2 tokenizer = Tokenizer(num_words=vocab_size)
3 tokenizer.fit_on_texts(x_train)
```

```
In [17]: 1 print('No. of reviews for training:',tokenizer.document_count)
2 #print(tokenizer.word_counts)
3 #print(tokenizer.word_docs)
4 #print(len(tokenizer.word_index))
```

No. of reviews for training: 112123

Convert data to integer sequence:

```
In [18]: 1 #CONVERT EACH DOC/REVIEW IN INTEGER SEQUENCE AND LENGTH OF ENCODED DOC=SIZE OF VOCAB
2 #x_tr = tokenizer.texts_to_matrix(x_train)
3 #CONVERT EACH DOC/REVIEW IN INTEGER SEQUENCE AND LENGTH OF ENCODED DOC=LEN OF DOC
4 x_train = tokenizer.texts_to_sequences(x_train)
5 x_test = tokenizer.texts_to_sequences(x_test)
```

```
In [19]: 1 print('sample of a review after encoding:\n',x_train[1])
```

sample of a review after encoding:

[13, 103, 164, 1659, 2, 1479, 1, 1029, 161, 93, 41, 92, 20, 4, 35, 1561, 3, 200, 286, 6, 6807, 7, 696, 95, 11, 1, 709, 3, 21, 54, 73, 176]

Padding the input data:


```
In [20]: 1 # truncate and/or pad input sequences
2 x_train = sequence.pad_sequences(x_train, maxlen=max_review_length)
3 x_test = sequence.pad_sequences(x_test, maxlen=max_review_length)
4
5 print('sample of a review after padding:\n',x_train[1])
```

sample of a review after padding:

[illegible]

0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	13	103	164	1659	2	1479
1	1029	161	93	41	92	20	4	35	1561	3	200	286	6
6807	7	696	95	11	1	709	3	21	54	73	176]		

Model-1:

[1.1]Model architecture:

```
In [21]: 1 model = Sequential()
2 model.add(Embedding(vocab_size, embedding_vecor_length, input_length=max_review_length))
3 model.add(LSTM(100, dropout=.5, recurrent_dropout=.5))
4 model.add(Dense(1, activation='sigmoid'))
5 print(model.summary())
```

WARNING:tensorflow:From /usr/local/lib/python3.5/dist-packages/tensorflow/python/framework/op_def_library.py:263: colocate_with (from tensorflow.python.framework.ops) is deprecated and will be removed in a future version.

Instructions for updating:

Colocations handled automatically by placer.

WARNING:tensorflow:From /usr/local/lib/python3.5/dist-packages/keras/backend/tensorflow_backend.py:3445: calling dropout (from tensorflow.python.ops.nn_ops) with keep_prob is deprecated and will be removed in a future version.

Instructions for updating:

Please use `rate` instead of `keep_prob`. Rate should be set to `rate = 1 - keep_prob`.

Layer (type)	Output Shape	Param #
=====		
embedding_1 (Embedding)	(None, 600, 32)	256000
=====		
lstm_1 (LSTM)	(None, 100)	53200
=====		
dense_1 (Dense)	(None, 1)	101
=====		
Total params: 309,301		
Trainable params: 309,301		
Non-trainable params: 0		
None		

[1.2]Train model on training data and draw loss plot:

```
In [24]: 1 %time model, history = trainModel(model, x_train, y_train, optimizer, loss, nb_epoch, batch_size)
2 saveModel(model, 'model[lstm(100)]')
```

Train on 78486 samples, validate on 33637 samples

Epoch 1/5

78486/78486 [=====] - 1556s 20ms/step - loss: 0.2148 - acc: 0.9172 - val_loss: 0.2110 - val_acc: 0.9212

Epoch 2/5

78486/78486 [=====] - 1561s 20ms/step - loss: 0.1976 - acc: 0.9231 - val_loss: 0.2089 - val_acc: 0.9227

Epoch 3/5

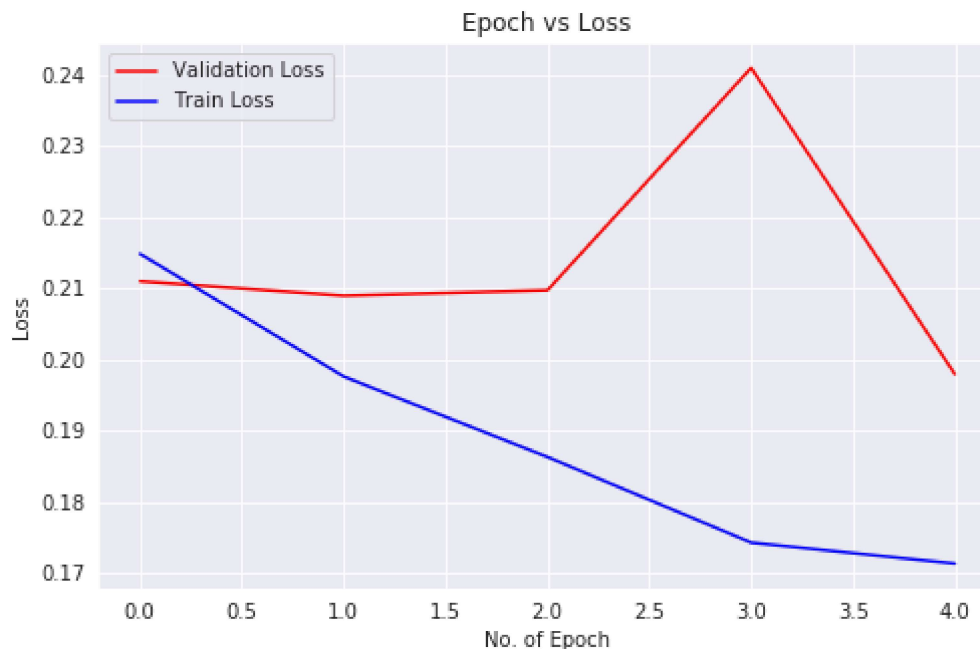
78486/78486 [=====] - 1549s 20ms/step - loss: 0.1863 - acc: 0.9281 - val_loss: 0.2097 - val_acc: 0.9217

Epoch 4/5

78486/78486 [=====] - 1559s 20ms/step - loss: 0.1743 - acc: 0.9329 - val_loss: 0.2409 - val_acc: 0.9000

Epoch 5/5

78486/78486 [=====] - 1569s 20ms/step - loss: 0.1713 - acc: 0.9342 - val_loss: 0.1979 - val_acc: 0.9258



CPU times: user 22h 4min 57s, sys: 1d 5h 43min 55s, total: 2d 3h 48min 53s

Wall time: 2h 9min 56s

[1.3] Model performance on test data:

```
In [28]: 1 local_summary=model_performance(x_test, y_test, nb_epoch, history, 'LSTM(100)')
```

```
48053/48053 [=====] - 362s 8ms/step
```

```
Test score: 0.1934
```

```
Test accuracy: 0.9260
```

[1.4]Model Summary:

```
In [29]: 1 print(local_summary)
```

```
+-----+-----+-----+-----+-----+-----+
| Model | Test-Loss | Test-Accuracy | Train-Loss | Train-Accuracy | Val-Loss | Val-Accu |
+-----+-----+-----+-----+-----+-----+
| LSTM(100) | 0.1934 | 0.9260 | 0.1713 | 0.9342 | 0.1979 | 0.9258 |
+-----+-----+-----+-----+-----+-----+
```

Model-2:**[2.1]Model architecture:**

```
In [32]: 1 #REFERENCE LINK: https://machinelearningmastery.com/return-sequences-and-return-states-for-lstms-in-keras/
2 model = Sequential()
3 model.add(Embedding(vocab_size, embedding_vector_length, input_length=max_review_length))
4 model.add(LSTM(100, return_sequences=True, dropout=.5, recurrent_dropout=.5))
5 model.add(LSTM(50, dropout=.5))
6 model.add(Dense(1, activation='sigmoid'))
7 model.summary()
```

Layer (type)	Output Shape	Param #
=====		
embedding_4 (Embedding)	(None, 600, 32)	256000
=====		
lstm_5 (LSTM)	(None, 600, 100)	53200
=====		
lstm_6 (LSTM)	(None, 50)	30200
=====		
dense_3 (Dense)	(None, 1)	51
=====		
Total params: 339,451		
Trainable params: 339,451		
Non-trainable params: 0		
=====		

[2.2]Train model on training data and draw loss plot:

```
In [35]: 1 %time model, history = trainModel(model, x_train, y_train, optimizer, loss, nb_epoch, batch_size)
        2 saveModel(model, 'model[lstm(100)-lstm(50)]')
```

Train on 78486 samples, validate on 33637 samples

Epoch 1/5

78486/78486 [=====] - 2774s 35ms/step - loss: 0.2898 - acc: 0.8870 - val_loss: 0.2124 - val_acc: 0.9178

Epoch 2/5

78486/78486 [=====] - 2779s 35ms/step - loss: 0.1987 - acc: 0.9223 - val_loss: 0.1980 - val_acc: 0.9241

Epoch 3/5

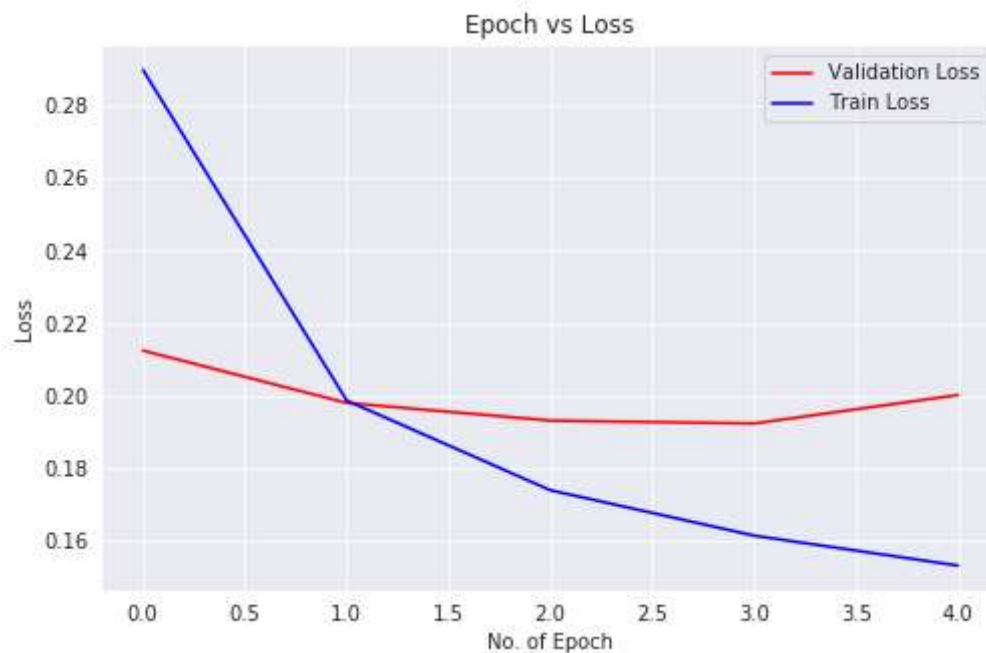
78486/78486 [=====] - 2780s 35ms/step - loss: 0.1740 - acc: 0.9332 - val_loss: 0.1932 - val_acc: 0.9246

Epoch 4/5

78486/78486 [=====] - 2806s 36ms/step - loss: 0.1614 - acc: 0.9380 - val_loss: 0.1923 - val_acc: 0.9248

Epoch 5/5

78486/78486 [=====] - 2804s 36ms/step - loss: 0.1532 - acc: 0.9414 - val_loss: 0.2002 - val_acc: 0.9277



CPU times: user 1d 15h 15min 5s, sys: 2d 5h 25min 57s, total: 3d 20h 41min 2s

Wall time: 3h 52min 25s

[2.3] Model performance on test data:

```
In [37]: 1 local_summary=model_performance(x_test, y_test, nb_epoch, history, 'LSTM(100)->LSTM(50)')
48053/48053 [=====] - 588s 12ms/step
Test score: 0.1962
Test accuracy: 0.9277
```

[2.4] Model Summary:

```
In [38]: 1 print(local_summary)
```

```
+-----+-----+-----+-----+-----+-----+
|      Model      | Test-Loss | Test-Accuracy | Train-Loss | Train-Accuracy | Val-Loss | Val-Accu |
+-----+-----+-----+-----+-----+-----+
| LSTM(100)->LSTM(50) |    0.1962 |    0.9277    |    0.1532 |    0.9414    |    0.2002 |    0.9277 |
+-----+-----+-----+-----+-----+-----+
```

Conclusion:

1. Performance with single LSTM with Dropout:

```
[a.] Loss : 0.1934
```

```
[b.] Acc  : 0.9260
```

2. Performance with stacked LSTM with Dropout:

```
[a.] Loss : 0.1962
```

```
[b.] Acc  : 0.9277
```

Got optimal performance with Stacked LSTM.

Reference Links:

1. <http://appliedaicourse.com/> (<http://appliedaicourse.com/>)
2. <https://machinelearningmastery.com/save-load-keras-deep-learning-models/> (<https://machinelearningmastery.com/save-load-keras-deep-learning-models/>)
3. <https://machinelearningmastery.com/stacked-long-short-term-memory-networks/> (<https://machinelearningmastery.com/stacked-long-short-term-memory-networks/>)
4. <https://machinelearningmastery.com/return-sequences-and-return-states-for-lstms-in-keras/> (<https://machinelearningmastery.com/return-sequences-and-return-states-for-lstms-in-keras/>)
5. <https://machinelearningmastery.com/prepare-text-data-deep-learning-keras/> (<https://machinelearningmastery.com/prepare-text-data-deep-learning-keras/>)